

Importance of regulariztion:

Preventing overfitting: Regularization helps to limit the model's capacity to memorize noise in the training data, leading to better generalization on unseen data.

Dealing with high-dimensional data: In datasets with a large number of features, regularization can help select the most important features and reduce the impact of irrelevant or redundant features.

Handling collinearity: When features are highly correlated, regularization techniques can mitigate the multicollinearity problem and provide more stable and interpretable coefficient estimates.

Improving model interpretability: Regularization can drive some coefficients to zero, effectively performing feature selection and simplifying the model, making it easier to interpret.

By incorporating regularization techniques into the model training process, we can achieve a balance between fitting the training data well and avoiding overfitting, leading to more robust and generalizable models.

Type of regulariztion:

1.Lasso regulariztion:convert high coefficient to 0 , this way eliminate unrelated attribute
 $L1_reg/Lasso R = loss + \alpha(|w|)$ where $\alpha(|w|)$ is penalty, α is parameter and w is the vector coefficient of model, regularization parameter control the strength of regularization

Ridge regression/L2 regularization: convert high coefficient to low coefficient of the model
 $L1_reg/Lasso R = loss + \alpha(|w|)^2$ where $\alpha(|w|)$ is penalty.

Elasticnet regression: $Elasticnet.R = loss + \alpha_1(|w|) + \alpha_2(|w|)^2$

$w = w_1 + w_2 + w_3 + \dots + w_n$ Noted, After regularization we trained the mode. In different model algo we can use Regularization as parameter l_1 and l_2 . for example NN, Logistic rregression SVM in here, we will use for genirilizing Lenear model

```
In [1]: #Import numerical libraries
import pandas as pd
import numpy as np

#Import graphical plotting libraries
import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline
```

```
In [6]: df = pd.read_csv('/Users/myyntiimac/Desktop/car-mpg.csv')
df.head()
```

```
Out[6]:
```

	mpg	cyl	disp	hp	wt	acc	yr	origin	car_type	car_name
0	18.0	8	307.0	130	3504	12.0	70	1	0	chevrolet chevelle malibu
1	15.0	8	350.0	165	3693	11.5	70	1	0	buick skylark 320
2	18.0	8	318.0	150	3436	11.0	70	1	0	plymouth satellite
3	16.0	8	304.0	150	3433	12.0	70	1	0	amc rebel sst
4	17.0	8	302.0	140	3449	10.5	70	1	0	ford torino

```
In [7]: df.shape
```

```
Out[7]: (398, 10)
```

Insight: This data contains 398 rows and 10 columns. By understanding 10 columns we can tell that miles per gallon (mpg) is dependent and the other nine variables are independent where we assume car_name has no effects in dependent variable mpg, so we can delete it by drop method.

Strating Basic EDA

```
In [8]: df = df.drop(['car_name'], axis = 1)
df.tail()
```

```
Out[8]:
```

	mpg	cyl	disp	hp	wt	acc	yr	origin	car_type
393	27.0	4	140.0	86	2790	15.6	82	1	1
394	44.0	4	97.0	52	2130	24.6	82	2	1
395	32.0	4	135.0	84	2295	11.6	82	1	1
396	28.0	4	120.0	79	2625	18.6	82	1	1
397	31.0	4	119.0	82	2720	19.4	82	1	1

```
In [9]: #Now check the null values in our df
df.isnull().any()
```

```
Out[9]: mpg          False
cyl            False
disp           False
hp             False
wt             False
acc            False
yr             False
origin         False
car_type       False
dtype: bool
```

```
In [11]: #check any column contain ? mark
df.isin(['?']).any()
```

```
Out[11]: mpg      False
          cyl      False
          disp     False
          hp       True
          wt       False
          acc      False
          yr       False
          origin   False
          car_type False
          dtype: bool
```

```
In [12]: #we find hp column contain ? mark
#now we replace the question mark with nan values then replace NAN velues wi
df = df.replace('?', np.nan)
```

```
In [13]: #check again
df.isin(['?']).any()
```

```
Out[13]: mpg      False
          cyl      False
          disp     False
          hp       False
          wt       False
          acc      False
          yr       False
          origin   False
          car_type False
          dtype: bool
```

```
In [14]: #NO ? mark now , but there is NAN values that need to fill
df['hp'].fillna(df['hp'].median(), inplace=True)
```

```
In [15]: #check again nan value
df.isnull().any()
```

```
Out[15]: mpg      False
          cyl      False
          disp     False
          hp       False
          wt       False
          acc      False
          yr       False
          origin   False
          car_type False
          dtype: bool
```

```
In [16]: #now no null value and ? mark in df
#But there is anouther column name "origin" which is catagorical column but
#this column we converted to individual column by get_dummies()
#For this first we create dictionary
df['origin'] = df['origin'].replace({1: 'america', 2: 'europe', 3: 'asia'})
df.head()
```

```
Out[16]:
```

	mpg	cyl	disp	hp	wt	acc	yr	origin	car_type
0	18.0	8	307.0	130	3504	12.0	70	america	0
1	15.0	8	350.0	165	3693	11.5	70	america	0
2	18.0	8	318.0	150	3436	11.0	70	america	0
3	16.0	8	304.0	150	3433	12.0	70	america	0
4	17.0	8	302.0	140	3449	10.5	70	america	0

```
In [17]: # now converted the origin column into 3 column which represent each contine
df = pd.get_dummies(df,columns = ['origin'])
```

```
In [18]: df.head()
```

```
Out[18]:
```

	mpg	cyl	disp	hp	wt	acc	yr	car_type	origin_america	origin_asia	origin_europe
0	18.0	8	307.0	130	3504	12.0	70	0	1	0	0
1	15.0	8	350.0	165	3693	11.5	70	0	1	0	0
2	18.0	8	318.0	150	3436	11.0	70	0	1	0	0
3	16.0	8	304.0	150	3433	12.0	70	0	1	0	0
4	17.0	8	302.0	140	3449	10.5	70	0	1	0	0

```
In [37]: df.shape
```

```
Out[37]: (398, 11)
```

feature scaling

Features often have different scales, ranges, or units of measurement. Scaling helps bring all the features to a similar scale, ensuring they have a similar impact during modeling. in our df , we can see the columns in different scale , so before training the model with training data we have to scalized, otherwise some featerure can show more effect on model(it will be bias)

Two scalization tecniques: 1) Z_score,Std.scaler/standerization=where mean of attribute convert to 0 and std.dev=1 $z = (x - \mu) / \sigma$

Where:

z is the standardized value of the data point. x is the original value of the data point. μ is the mean of the feature. σ is the standard deviation of the feature. 2)Normalization /min_max scaler:values are shifted 0 to 1 $x_{scaled} = (x - \min(x)) / (\max(x) - \min(x))$

Where:

x is the original value of the data point. x_{scaled} is the scaled value of the data point. $\min(x)$ is the minimum value of the feature. $\max(x)$ is the maximum value of the feature.

First we divide the data into independent (X) and dependent data (y) then we scale it. WHY?

Because *The reason we don't scale the entire data before and then divide it into train(X) & test(y) is because once you scale the data, the type(data_s) would be numpy.ndarray. It's impossible to divide this data when it's an array.*

Hence we divide type(data) pandas.DataFrame, then proceed to scaling it.

```
In [38]: x = df.drop(['mpg'], axis = 1) # independent variable
        y = df[['mpg']]
```

```
In [39]: x
```

```
Out[39]:
```

	cyl	displacement	horsepower	weight	acceleration	year	car_type	origin_america	origin_asia	origin_europe
0	8	307.0	130	3504	12.0	70	0	1	0	0
1	8	350.0	165	3693	11.5	70	0	1	0	0
2	8	318.0	150	3436	11.0	70	0	1	0	0
3	8	304.0	150	3433	12.0	70	0	1	0	0
4	8	302.0	140	3449	10.5	70	0	1	0	0
...
393	4	140.0	86	2790	15.6	82	1	1	0	0
394	4	97.0	52	2130	24.6	82	1	0	0	1
395	4	135.0	84	2295	11.6	82	1	1	0	0
396	4	120.0	79	2625	18.6	82	1	1	0	0
397	4	119.0	82	2720	19.4	82	1	1	0	0

398 rows × 10 columns

```
In [40]: y
```

```
Out[40]:
```

	mpg
0	18.0
1	15.0
2	18.0
3	16.0
4	17.0
...	...
393	27.0
394	44.0
395	32.0
396	28.0
397	31.0

398 rows × 1 columns

```
In [46]: #Now we scalize the independent and dependent variable by std.scalirization
        from sklearn.preprocessing import StandardScaler

        # Standardize the independent variables
        scaler = StandardScaler()
        X_std = scaler.fit_transform(X)
        X_std = pd.DataFrame(X_std, columns = X.columns) #converting scaled data into
        X_std
```

Out [46]:

	cyl	disp	hp	wt	acc	yr	car_type	origin_an
0	1.498191	1.090604	0.673118	0.630870	-1.295498	-1.627426	-1.062235	0.7
1	1.498191	1.503514	1.589958	0.854333	-1.477038	-1.627426	-1.062235	0.7
2	1.498191	1.196232	1.197027	0.550470	-1.658577	-1.627426	-1.062235	0.7
3	1.498191	1.061796	1.197027	0.546923	-1.295498	-1.627426	-1.062235	0.7
4	1.498191	1.042591	0.935072	0.565841	-1.840117	-1.627426	-1.062235	0.7
...
393	-0.856321	-0.513026	-0.479482	-0.213324	0.011586	1.621983	0.941412	0.7
394	-0.856321	-0.925936	-1.370127	-0.993671	3.279296	1.621983	0.941412	-1.2
395	-0.856321	-0.561039	-0.531873	-0.798585	-1.440730	1.621983	0.941412	0.7
396	-0.856321	-0.705077	-0.662850	-0.408411	1.100822	1.621983	0.941412	0.7
397	-0.856321	-0.714680	-0.584264	-0.296088	1.391285	1.621983	0.941412	0.7

398 rows × 10 columns

In [48]:

```
scaler = StandardScaler()
y_std = scaler.fit_transform(y)
y_std = pd.DataFrame(y_std, columns = y.columns) #converting scaled data into y_std
```

Out [48]:

	mpg
0	-0.706439
1	-1.090751
2	-0.706439
3	-0.962647
4	-0.834543
...	...
393	0.446497
394	2.624265
395	1.087017
396	0.574601
397	0.958913

398 rows × 1 columns

Insight:The values in column 0 have been transformed such that the mean is 0 and the standard deviation is 1. The first value (-0.856) is below the mean and the other value (1.498) is above the mean.

```
In [49]: #now we split the both variable for training data and test data
# for this we import test_train_split() from skleran_model_selection
from sklearn.model_selection import train_test_split
X_train, X_test, y_train,y_test = train_test_split(X_std, y_std, test_size =
X_train.shape
```

```
Out[49]: (298, 10)
```

```
In [ ]: #Now build the model
#First Lenear model then lasso and ridge and compare the performance of the
# FOr first lenear model we use enamurate() to see how model coefficient loc
#and decrease the overfitting problem and improve the model
```

```
In [50]: from sklearn.linear_model import LinearRegression

# Create and fit the linear regression model
regression_model = LinearRegression()
regression_model.fit(X_train,y_train)

# Print the coefficients for each independent variable
#oefficients for each independent variable are printed using a for loop comb
for col_name, coef in zip(X_train.columns, regression_model.coef_[0]):
    print(f"The coefficient for {col_name} is {coef}")

# Print the intercept
#The intercept is stored in the intercept variable and printed using an f-st
intercept = regression_model.intercept_[0]
print(f"The intercept is {intercept}")
```

```
The coefficient for cyl is 0.295188274518623
The coefficient for disp is 0.34377235544447493
The coefficient for hp is -0.2031300205066636
The coefficient for wt is -0.7282388801097291
The coefficient for acc is 0.03466225460648925
The coefficient for yr is 0.3805726133418513
The coefficient for car_type is 0.36210864572959445
The coefficient for origin_america is -0.08228452280158736
The coefficient for origin_asia is 0.05237035868641345
The coefficient for origin_europe is 0.04973173878319417
The intercept is 0.021346953567712677
```

```
In [52]: #Ridge regression
#alpha factor here is lambda (penalty term) which helps to reduce the magnit
from sklearn.linear_model import Ridge
ridge_model = Ridge(alpha = 0.3)
ridge_model.fit(X_train, y_train)

print('Ridge model coef: {}'.format(ridge_model.coef_))
```

```
Ridge model coef: [[ 0.292303    0.33174595 -0.20338688 -0.7180704    0.03282
392  0.37944512
 0.3587562  -0.08172261  0.05226426  0.04912861]]
```

Insight:we find that coefficient is changed

```
In [53]: #Lets check in Lasso regression
from sklearn.linear_model import Lasso
Lasso_model = Lasso(alpha = 0.1)
Lasso_model.fit(X_train, y_train)

print('Lasso model coef: {}'.format(Lasso_model.coef_))
```

```
Lasso model coef: [-0.          -0.          -0.02288344 -0.52067427  0.
0.2890683
0.11160748 -0.02891466  0.          0.          ]
```

Insight: we found some coegffient are 0, thats way lasso , eleminate or ignore the complex coefficient

```
In [55]: # now compare the model with R square value
#The score() method returns the R-squared value,
#Simple Linear Model
print(regression_model.score(X_train, y_train))
print(regression_model.score(X_test, y_test))

print('*****')
#Ridge
print(ridge_model.score(X_train, y_train))
print(ridge_model.score(X_test, y_test))

print('*****')
#Lasso
print(Lasso_model.score(X_train, y_train))
print(Lasso_model.score(X_test, y_test))
```

```
0.8343520392348843
0.8575776228871523
*****
0.8343385446094193
0.8582584230696331
*****
0.7952115294892155
0.854782394138887
```

In []: Insight:Both Ridge & Lasso regularization performs very well on this data, t

```
In [ ]: #Model parameter tuning
#r^2 is not a reliable metric as it always increases with addition of more a
#Instead we use adjusted r^2 which removes the statistical chance that impro
#Scikit does not provide a facility for adjusted r^2...
#so we use statsmodel, a library that gives results similar to what you obta
```

```
In [58]: import statsmodels.api as sm
# Prepare the data
X = X_train[['cyl', 'disp', 'hp', 'wt', 'acc', 'yr', 'car_type', 'origin_ame
y = y_train['mpg']

# Add constant term to the independent variables
X = sm.add_constant(X)

# Fit the OLS model
ols_model = sm.OLS(y, X).fit()

# Get the parameter estimates
params = ols_model.params
params
```



```
Out[58]: const      0.021347
          cyl        0.295188
          disp       0.343772
          hp         -0.203130
          wt         -0.728239
          acc        0.034662
          yr         0.380573
          car_type    0.362109
          origin_america -0.082285
          origin_europe  0.049732
          origin_asia   0.052370
          dtype: float64
```

```
In [60]: ols_model.summary()
```

Out [60]:

OLS Regression Results							
Dep. Variable:		mpg		R-squared:		0.834	
Model:		OLS		Adj. R-squared:		0.829	
Method:		Least Squares		F-statistic:		161.2	
Date:		Tue, 13 Jun 2023		Prob (F-statistic):		6.34e-107	
Time:		03:54:43		Log-Likelihood:		-159.88	
No. Observations:		298		AIC:		339.8	
Df Residuals:		288		BIC:		376.7	
Df Model:		9					
Covariance Type:		nonrobust					
		coef	std err	t	P> t	[0.025	0.975]
const	0.0213	0.024	0.872	0.384	-0.027	0.070	
cyl	0.2952	0.106	2.776	0.006	0.086	0.504	
disp	0.3438	0.124	2.779	0.006	0.100	0.587	
hp	-0.2031	0.076	-2.680	0.008	-0.352	-0.054	
wt	-0.7282	0.086	-8.476	0.000	-0.897	-0.559	
acc	0.0347	0.039	0.900	0.369	-0.041	0.110	
yr	0.3806	0.028	13.454	0.000	0.325	0.436	
car_type	0.3621	0.065	5.586	0.000	0.235	0.490	
origin_america	-0.0823	0.020	-4.202	0.000	-0.121	-0.044	
origin_europe	0.0497	0.021	2.392	0.017	0.009	0.091	
origin_asia	0.0524	0.020	2.647	0.009	0.013	0.091	
Omnibus:		17.430	Durbin-Watson:		2.142		
Prob(Omnibus):		0.000	Jarque-Bera (JB):		23.858		
Skew:		0.439	Prob(JB):		6.60e-06		
Kurtosis:		4.072	Cond. No.		6.69e+15		

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The smallest eigenvalue is 3.75e-29. This might indicate that there are strong multicollinearity problems or that the design matrix is singular.

```
In [61]: # calculate MSE from predicted Y and actual Y_test
Y_pred=regression_model.predict(X_test)
Y_pred
```

```
Out[61]: array([[-5.45633545e-01],
 [ 6.05224257e-01],
 [-2.95500865e-01],
 [ 6.02203961e-01],
 [-9.01910406e-02],
 [-8.37764197e-01],
 [ 8.67029488e-01],
 [ 1.50701925e+00],
 [-6.28044713e-01],
 [-1.58460998e+00],
 [ 9.18187719e-01],
 [-6.37225904e-01],
 [-4.30911078e-01],
 [ 4.14159104e-01],
 [ 1.74426214e+00],
 [-3.75431930e-03],
 [-1.62418193e+00],
 [-6.25345918e-01],
 [-1.82223398e+00],
 [ 1.32148818e+00],
 [ 3.71511216e-01],
 [ 1.09488132e+00],
 [-5.40607558e-01],
 [ 2.60428937e-01],
 [ 3.72999854e-01],
 [ 9.26136877e-01],
 [ 1.23596475e+00],
 [ 1.29715456e+00],
 [-9.60468185e-01],
 [ 8.74402695e-01],
 [ 1.97066475e-01],
 [-1.68249676e+00],
 [-2.24031935e-01],
 [ 6.92675135e-01],
 [ 2.49334561e-01],
 [-1.20816136e+00],
 [ 4.58206904e-01],
 [-1.88550614e+00],
 [ 1.06084984e+00],
 [ 1.56106357e-01],
 [ 1.41286332e-01],
 [ 1.92268928e-01],
 [-1.75606207e-01],
 [ 1.33575371e+00],
 [ 1.79737791e-03],
 [-4.05584061e-01],
 [-4.93245447e-01],
 [-1.42040779e+00],
 [ 7.25390608e-01],
 [-7.98713325e-01],
 [ 1.57375909e-01],
 [ 4.12742795e-01],
 [-7.13208661e-01],
 [-1.34085256e+00],
 [ 6.72307620e-01],
 [ 2.37617863e-01],
 [-1.71012208e+00],
 [-1.21581935e+00],
 [ 9.83867271e-01],
 [ 1.63431074e+00],
 [ 1.63346651e+00],
 [ 1.64312033e+00],
 [-6.01732409e-01],
 [ 2.09207571e-01],
```

```

[-2.86040505e-01],
[ 1.24485074e+00],
[ 7.20692290e-02],
[ 1.83589350e-01],
[ 8.41495172e-01],
[-1.27985279e+00],
[-2.62295832e-01],
[-5.39652159e-02],
[-1.11067555e+00],
[ 4.46663210e-01],
[-1.38343143e+00],
[ 3.18637749e-01],
[ 9.13874982e-01],
[-7.67583552e-01],
[-1.32838570e+00],
[-1.18141950e-02],
[-9.62087107e-02],
[-6.01430932e-01],
[ 1.61760033e+00],
[ 9.78412826e-02],
[ 1.50019976e+00],
[-9.66352149e-01],
[-8.58942279e-01],
[-1.58472597e-01],
[ 1.55463479e+00],
[ 1.12368955e+00],
[-3.54978399e-01],
[-3.72939663e-01],
[-1.83544510e-01],
[ 5.36823685e-01],
[ 1.01009679e+00],
[-3.92142556e-01],
[ 1.33078142e+00],
[ 6.84414144e-01],
[ 6.77705360e-02],
[-7.29293202e-02]])

```

```

In [63]: # Calculate Mean Squared Error (MSE)
mse = np.mean((Y_pred - y_test) ** 2)
mse

```

```

/Users/myyntiimac/anaconda3/lib/python3.10/site-packages/numpy/core/fromnumeric.py:3430: FutureWarning: In a future version, DataFrame.mean(axis=None) will return a scalar mean over the entire DataFrame. To retain the old behavior, use 'frame.mean(axis=0)' or just 'frame.mean()'
  return mean(axis=axis, dtype=dtype, out=out, **kwargs)

```

```

Out[63]: mpg      0.128032
dtype: float64

```

```

In [64]: # calculate Root Mean square error
import math
rmse = math.sqrt(mse)
rmse

```

```

Out[64]: 0.35781548331897894

```

Insight: so the difference between actual mpg and predicted mpg is .357

Chceck the quality of regression model

Residual plots are useful for evaluating the assumptions of a regression model.

They show the difference between the observed values and the predicted values (i.e., the residuals) on the y-axis,

while the x-axis represents the independent variable values. The lowess curve in the plot provides a smoothed representation of the residuals,


which can help identify patterns or nonlinear relationships between variables.


```
In [66]: #Lets check the residuals for some of these predictor.

fig = plt.figure(figsize=(10,10))
sns.residplot(x= X_test['hp'], y= y_test['mpg'], color='red', lowess=True )

fig = plt.figure(figsize=(10,8))
sns.residplot(x= X_test['yr'], y= y_test['mpg'], color='black', lowess=True

Out[66]: <Axes: xlabel='yr', ylabel='mpg'>
```






The x-axis represents the year of the car(independent variable). The y-axis represents the residuals, which are the differences between the actual car mpg and the predicted mpg. Each point on the scatterplot indicates how well the model's predictions match the actual mpg. If the points are randomly scattered around the horizontal line at $y=0$, it suggests that the model's predictions are unbiased and accurate. Horizontal Line at $y=0$:

The black line at $y=0$ indicates the reference line, representing zero residual. It helps to visually assess whether the residuals are centered around zero. If the points are evenly distributed above and below the line, it indicates that the model is making unbiased predictions on average.

```
In [68]: plt.scatter(y_test['mpg'], Y_pred)
```

```
Out[68]: <matplotlib.collections.PathCollection at 0x7f8502f15bd0>
```



Insight:corelation between predicted and actual mpg is positive